

## UNIT-I

### 1.1 Database-System Applications

Explains where database systems are used in real life:

- Banking (accounts, loans)
  - Universities (student records, courses)
  - Airlines (reservations)
  - E-commerce, human resources, telecommunications, social media, scientific data, etc.
- Emphasizes that modern applications need to manage large, persistent, shared, and reliable data.

### 1.2 Purpose of Database Systems

Contrasts the traditional **file-processing approach** (data redundancy, inconsistency, difficult concurrency, weak security) with **database systems**: Advantages:

- Data abstraction and independence
- Reduced redundancy
- Consistency and integrity constraints
- Efficient querying
- Concurrent access and crash recovery
- Security and authorization

### 1.3 View of Data

Introduces levels of abstraction:

1. **Physical level** (how data is stored on disk)
2. **Logical level** (schema: tables, columns, relationships)
3. **View level** (external schemas or views for different users)

Also covers **data models**:

- Relational model
- Entity-Relationship (E-R) model
- Object-based, semi-structured (XML), NoSQL, etc.

### 1.4 Database Languages

- **Data-Definition Language (DDL)** → defines schema (CREATE TABLE, etc.)
- **Data-Manipulation Language (DML)** → querying and updates
  - Procedural (e.g., relational algebra)
  - Declarative (e.g., SQL: SELECT, INSERT, UPDATE, DELETE)
- Storage definition language (SDL), view definition language (VDL)

- Modern systems combine DDL and DML in SQL

(Exam-Ready Summary – Silberschatz Style)

Language Type	Full Name	Purpose	Examples (SQL)
<b>DDL</b>	Data Definition Language	Define, alter, or drop schema (structure) of database objects	CREATE TABLE, ALTER TABLE, DROP TABLE, CREATE INDEX
<b>DML</b>	Data Manipulation Language	Query and modify data (insert, update, delete, retrieve)	SELECT, INSERT, UPDATE, DELETE
<b>DCL</b>	Data Control Language	Security: grant/revoke permissions	GRANT, REVOKE
<b>TCL</b>	Transaction Control Language	Manage transactions	COMMIT, ROLLBACK, SAVEPOINT
<b>SDL / VDL</b>	Storage / View Definition Language	Rarely used separately in modern DBMS (merged into DDL)	(Historical – now part of SQL DDL)

SQL

-- *DDL*

```
CREATE TABLE Student (
    id      INT PRIMARY KEY,
    name    VARCHAR(50) NOT NULL,
    dept    VARCHAR(10)
);
```

-- *DML*

```
INSERT INTO Student VALUES (101, 'Ali', 'CS');
UPDATE Student SET dept = 'EE' WHERE id = 101;
DELETE FROM Student WHERE id = 101;
SELECT * FROM Student WHERE dept = 'CS';
```

-- *View (part of DDL)*

```
CREATE VIEW CS_Students AS
SELECT id, name FROM Student WHERE dept = 'CS';
```

-- DCL

```
GRANT SELECT ON Student TO public;  
REVOKE INSERT ON Student FROM user_bob;
```

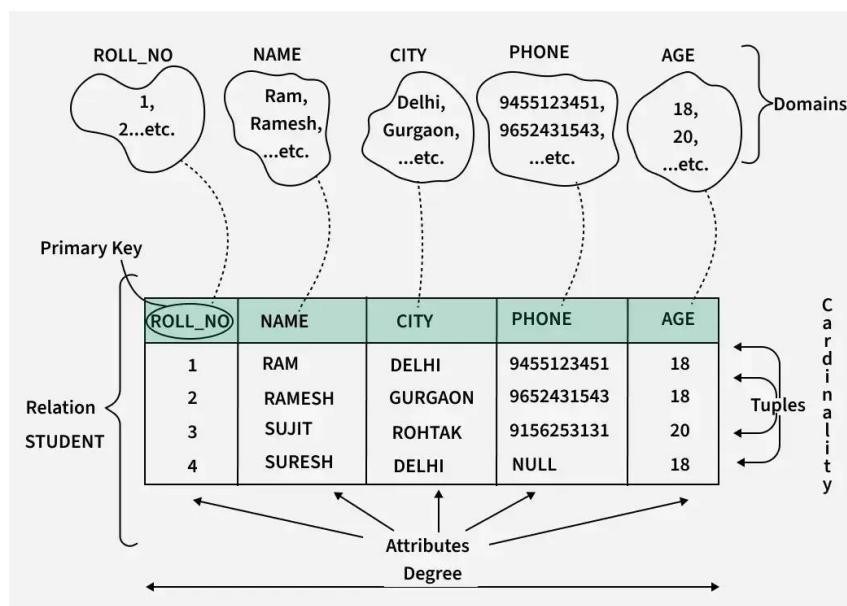
-- TCL

```
BEGIN TRANSACTION;  
UPDATE Account SET balance = balance - 100 WHERE id = 1;  
UPDATE Account SET balance = balance + 100 WHERE id = 2;  
COMMIT; -- or ROLLBACK if something went wrong
```

## 1.5 Relational Databases

Core concepts:

- Relation = table
- Tuple = row
- Attribute = column
- Domain, keys (primary, candidate, foreign)
- Relational schema and instance
- Brief introduction to relational algebra and SQL as the standard language



## Types of Keys in the Relational Model

**Primary Key:** A Primary Key uniquely identifies each tuple in a relation. It must contain unique values and cannot have NULL values.

**Candidate Key:** A Candidate Key is a set of attributes that can uniquely identify a tuple in a relation.

**Super Key:** A Super Key is a set of attributes that can uniquely identify a tuple.

**Foreign Key:** A Foreign Key is an attribute in one relation that refers to the primary key of another relation.

**Composite Key:** A Composite Key is formed by combining two or more attributes to uniquely identify a tuple.

## 1.6 Database Design

Design process:

1. Requirement analysis
2. Conceptual design using **E-R model** (entities, attributes, relationships, cardinality, participation, weak entities, specialization/generalization)
3. Logical design → mapping E-R to relational model
4. Schema refinement and normalization (1NF to BCNF, functional dependencies)
5. Physical design (indexes, file organization)

## 1.7 Data Storage and Querying

Storage hierarchy:

- Disk → buffer manager → files and access methods → indexes (B+-tree, hash)

Query processing:

- Parsing → query rewrite → optimization → evaluation plan → execution Cost-based optimization using statistics

## 1.8 Transaction Management

Transaction = unit of work that must be **atomic, consistent, isolated, durable (ACID)**

- Concurrency control (locks, timestamp ordering, multiversion)
- Recovery (log-based, shadow paging, ARIES)
- Deadlock handling

## ACID Properties (Must Memorize Exactly)

Property	Meaning	How DBMS Guarantees It
Atomicity	All operations succeed or none (all-or-nothing)	Logging + Undo (ROLLBACK)
Consistency	Takes DB from one consistent state to another	Constraints + Application logic + Atomicity
Isolation	Partial effects of incomplete transactions invisible to others	Concurrency control (locks, MVCC, timestamps)
Durability	Once committed, changes survive crashes	Write-Ahead Logging (WAL) + Force-write to disk

## Concurrency Control (How multiple transactions run safely)

Technique	How it Works	Pros / Cons
Locking (2PL)	Shared (S) lock for read, Exclusive (X) lock for write	Simple, strict 2PL guarantees serializability
	Strict 2PL: hold X locks until commit	Prevents cascading rollback
Timestamp Ordering	Each transaction gets timestamp; enforce order on conflicting operations	Deadlock-free, but can abort many times
Multiversion CC	Keeps old versions → readers never block writers (used in PostgreSQL, Oracle)	High concurrency, no reader blocking
Validation (Optimistic)	Read → Work → Validate (no conflict?) → Commit	Good when conflicts rare

## 1.9 Database Architecture

- Centralized vs. client-server vs. parallel vs. distributed databases
- Three-tier architecture (presentation, application/logic, database)
- Data servers, transaction servers

### Database Architecture

Architecture Type	Description	Typical Use-Cases / Examples (2025)
Centralized	Single computer, single DBMS instance	Small university department, legacy systems

Architecture Type	Description	Typical Use-Cases / Examples (2025)
Client–Server	Clients (apps) connect via network to a dedicated database server	Most traditional web apps (PHP + MySQL, Java + Oracle)
Two-Tier	Client machine runs application + directly talks to DB server	Desktop apps, early web apps
Three-Tier (most common today)	1. Presentation tier (browser/mobile app) 2. Application/logic tier (Node.js, Django, Spring, etc.) 3. Database tier	Netflix, Instagram, Uber, Amazon, almost every modern web/mobile app
Parallel Databases	Multiple CPUs/disks in one machine or cluster – shared-memory, shared-disk, shared-nothing	Google BigQuery, Snowflake, Teradata, Vertica
Shared-Nothing (winner for scale)	Each node has its own CPU + memory + disk; only network connects them	YouTube (Vitess), TikTok, Facebook (MyRocks), Uber (Schemaless)
Distributed Databases	Data spread across geographically different sites, still appears as one logical DB	Google Spanner, CockroachDB, YugabyteDB, Planet-scale (Vitess)
Cloud-Native Databases	Fully managed, auto-scaling, multi-region replication	AWS Aurora, DynamoDB Global Tables, Azure Cosmos DB, Neon (serverless Postgres)

## 1.10 Data Mining and Information Retrieval

- **Data mining:** discovering patterns (classification, clustering, association rules)
- **Information retrieval:** text search, ranking (e.g., PageRank), handling unstructured data  
Briefly distinguishes traditional DBMS from these newer areas

## 1.12 Database Users and Administrators

User categories:

- End users (naive, sophisticated)
- Application programmers
- Database administrator (DBA): schema definition, storage and authorization, backup/recovery, performance tuning, security

## Database Users and Administrators

User Type	Who They Are	How They Interact with the DBMS	Real-Life Examples (2025)
<b>Naive / Casual Users</b>	Non-technical end users, no SQL knowledge	Use mobile/web apps or ATMs – never write queries	You checking bank balance on PhonePe Passenger booking Uber Student viewing grades on portal
<b>Sophisticated Users</b>	Analysts, engineers, scientists	Write complex SQL / analytical queries, use BI tools	Data analyst at Amazon writing “top 100 products this month” Researcher querying genomic database
<b>Application Programmers</b>	Software developers	Write code in Python/Java/Go/Node that embeds SQL or ORM	Backend engineer at Swiggy writing INSERT INTO orders ... Full-stack dev using Django ORM
<b>Database Administrators (DBA)</b>	The “god” of the database	Highest privileges – everything below + system-level tasks	Person who upgrades PostgreSQL at 3 AM, creates backups, tunes indexes, grants permissions

Section	Concept	Real-World Example
1.1 Database-System Applications	Where databases are used	<ul style="list-style-type: none"> <li>Spotify: stores billions of songs, user playlists, listening history</li> <li>Uber: matches riders and drivers in real time using location data</li> <li>Amazon: product catalog, orders, reviews, recommendations</li> <li>Hospital: patient records, appointments, prescriptions</li> </ul>
1.2 Purpose of Database Systems	Problems of file-processing vs. DBMS	Without DBMS (old way): HR stores employee data in Excel, Payroll stores it in another Excel → same employee has different salaries → inconsistency! With DBMS (modern way): Single Oracle/MySQL database → one source of truth, automatic consistency
1.3 View of Data	Three levels of abstraction	University database: <ul style="list-style-type: none"> <li>Physical level: data stored in B+-tree indexes on SSD</li> <li>Logical level: tables Student(ID, name, dept), Enrolled(sid, cid, grade)</li> <li>View level:               <ul style="list-style-type: none"> <li>Student sees only own grades: CREATE VIEW MyGrades AS SELECT ...</li> <li>Registrar sees everything</li> <li>Professor sees only own classes</li> </ul> </li> </ul>

Section	Concept	Real-World Example
1.4 Database Languages	DDL vs DML (SQL)	DDL (schema): CREATE TABLE Student (id INT PRIMARY KEY, name VARCHAR(50)); DML (data): INSERT INTO Student VALUES (12345, 'Ali'); SELECT name FROM Student WHERE id = 12345;
1.5 Relational Databases	Tables, keys, foreign keys	Banking example: Customer(customer_id PK, name, address) Account(account_no PK, balance, customer_id FK → Customer)
1.6 Database Design	E-R → Relational (simplified)	Library E-R: Entity: Book(ISBN, title), Borrower(card_no, name) Relationship: Borrows (many-to-many) → Relational tables: Book(ISBN PK, title) Borrower(card_no PK, name) Borrows(ISBN FK, card_no FK, borrow_date, due_date)
1.7 Data Storage and Querying	Indexes speeding up queries	Amazon searching 300 million products by name: Without index → 5–10 seconds With B+-tree index on product_name → < 50 ms
1.8 Transaction Management	ACID in action	ATM withdrawal of \$100: 1. Check balance $\geq$ 100 2. Deduct \$100 from account 3. Dispense cash If power fails after step 2, recovery system rolls back → money is not lost (Durability & Atomicity)
1.9 Database Architecture	Three-tier architecture	Instagram (mobile app): • Tier 1 (Presentation): Your phone screen • Tier 2 (Application): Python/Django servers that apply filters, generate news feed • Tier 3 (Database): Thousands of MySQL + Cassandra shards storing posts, likes, comments
1.10 Data Mining and Information Retrieval	Mining vs Retrieval	• Data mining: Netflix discovers “people who watched Stranger Things also watched Wednesday” → recommendation • Information retrieval: You type “best noise-cancelling headphones” in Google → ranks pages using PageRank + relevance
1.12 Database Users and Administrators	Different users	• Naive user: You checking your bank balance on the mobile app • Sophisticated user: Data analyst writing complex SQL to find “top 10 customers by spending this month” • Application programmer: Developer writing the Python backend for the banking app • DBA: Person who upgrades PostgreSQL at 3 a.m. when no one is using the system

Company / App	What Happens in Real Time	Main Database(s) Used	Why This DB? / Special Techniques
Uber	Rider opens app → sees cars moving, gets ETA, live pricing	<ul style="list-style-type: none"> <li>• <b>Google Spanner</b> (global rides)</li> <li>• <b>Schemaless (MySQL fork)</b></li> <li>• <b>Redis</b> (real-time cache)</li> </ul>	Spanner = globally consistent in milliseconds Redis for driver location (billions of writes/sec)
WhatsApp	100+ billion messages per day, all delivered instantly	<ul style="list-style-type: none"> <li>• <b>SQLite</b> (on phone)</li> <li>• <b>Custom Erlang-based DB</b> (backend)</li> <li>• <b> RocksDB</b> for message storage</li> </ul>	Extremely high write throughput, low latency
TikTok	You scroll → infinite personalized video feed in <200 ms	<ul style="list-style-type: none"> <li>• <b>TiDB + MySQL</b> (user data)</li> <li>• <b>Redis + SSDB</b> (feed cache)</li> <li>• <b>ClickHouse</b> (analytics)</li> </ul>	Recommendation engine queries DB 1000s times per scroll
Google Maps	1 billion users see live traffic, rerouting in seconds	<ul style="list-style-type: none"> <li>• <b>Spanner</b> (core map data)</li> <li>• <b>Bigtable</b> (traffic layers)</li> <li>• <b>Memorystore (Redis)</b></li> </ul>	Spanner handles planet-scale consistent reads/writes
Stock Exchanges (NYSE, Binance, Nasdaq)	Millions of trades per second, matching buy/sell instantly	<ul style="list-style-type: none"> <li>• <b>kdb+</b> (primary for high-frequency trading)</li> <li>• <b>Aerospike / Redis</b> (order book)</li> <li>• <b>PostgreSQL</b> (settlement)</li> </ul>	kdb+ can do millions of ops/sec in memory
Airbnb	You search “Paris, 3 guests, Dec 24–27” → instant results with live availability	<ul style="list-style-type: none"> <li>• <b>MySQL + Vitess</b> (sharded)</li> <li>• <b>Redis</b> (search cache)</li> <li>• <b>Elasticsearch</b> (full-text)</li> </ul>	Vitess scales MySQL to millions of QPS
Netflix	200+ million users streaming, recommendations update live	<ul style="list-style-type: none"> <li>• <b>Cassandra</b> (user profiles, viewing history)</li> <li>• <b>DynamoDB</b> (some regions)</li> <li>• <b>Redis</b> (session &amp; queue)</li> </ul>	Cassandra handles massive write scaling

Company / App	What Happens in Real Time	Main Database(s) Used	Why This DB? / Special Techniques
Swiggy / DoorDash / Zomato	You place food order → restaurant + delivery partner get it instantly	<ul style="list-style-type: none"> <li>• PostgreSQL (main data)</li> <li>Redis (real-time order tracking)</li> <li>• Kafka + Flink (event streaming)</li> </ul>	Real-time ETA calculation using live GPS + Redis
Paytm / PhonePe (India)	UPI payments settle in <1 second, 10M+ transactions/min	<ul style="list-style-type: none"> <li>• Custom sharded MySQL</li> <li>• Redis (fraud detection)</li> <li>• Aerospike (real-time balance)</li> </ul>	Sub-100ms latency required by government regulation
Discord	150M monthly users chatting, voice, live screen sharing	<ul style="list-style-type: none"> <li>• Cassandra (chat history)</li> <li>• ScyllaDB (newer clusters)</li> <li>• Redis (online presence)</li> </ul>	ScyllaDB = drop-in Cassandra replacement, 5–10× faster
Tesla (in-car system)	Live navigation, Autopilot data upload, over-the-air updates	<ul style="list-style-type: none"> <li>• SQLite (per car)</li> <li>• Cassandra + Kafka (fleet data)</li> <li>• ClickHouse (telemetry analytics)</li> </ul>	Millions of cars uploading sensor data every minute
Formula 1 Timing	20 cars, 300+ sensors each, data shown live on TV in <0.1s	<ul style="list-style-type: none"> <li>• Oracle Database + Berkeley DB (real-time)</li> <li>• Custom C++ in-memory engine</li> </ul>	Official timing partner uses Oracle for decades